

HIT391

MACHINE LEARNING: ADVANCEMENTS AND APPLICATIONS



Week 7:

Support Vector Machines

- **Learning Outcomes**
 - Hyperplane
 - Maximal Margin Classifier
 - Support Vector Classifier
 - Feature Expansion

Support Vector Machines

- Here we approach the two-class classification problem in a direct way:
 - We try and find a plane that **separates the classes** in **feature space**.
- **What if we can't linearly separate the data?**

When a simple line (or hyperplane) can't separate the data, we get creative in two main ways:

1. Support Vector Classifiers (Soft Margin)

- Instead of requiring a perfect separation, we **allow some misclassifications**.
- This gives us a “**soft margin**”—a more flexible boundary that can handle overlap or noise in the data.

2. Data Transformations (Feature Expansion)

- We transform the data into a new space where separation is possible.
- This is like lifting the data into a higher-dimensional space where complex patterns **become linear**.

Hyperplane

- A hyperplane in p dimensions is a flat affine subspace of dimension $p - 1$.
- In general the equation for a **hyperplane** in p -dimensional space has the form:

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$

- when the dimension $p = 2$, the hyperplane is a **line (1D)**.
- when $\beta_0 = 0$, the hyperplane goes through **the origin**, otherwise not.
- The vector $\beta = (\beta_1, \beta_2, \dots, \beta_p)$ is called the **normal vector** — it points in a **direction orthogonal to the surface of a hyperplane**.

An example ($p=2$)

1. **The hyperplane(blue)** – separating the space.

2. **The normal vector β (red)** – controls orientation.

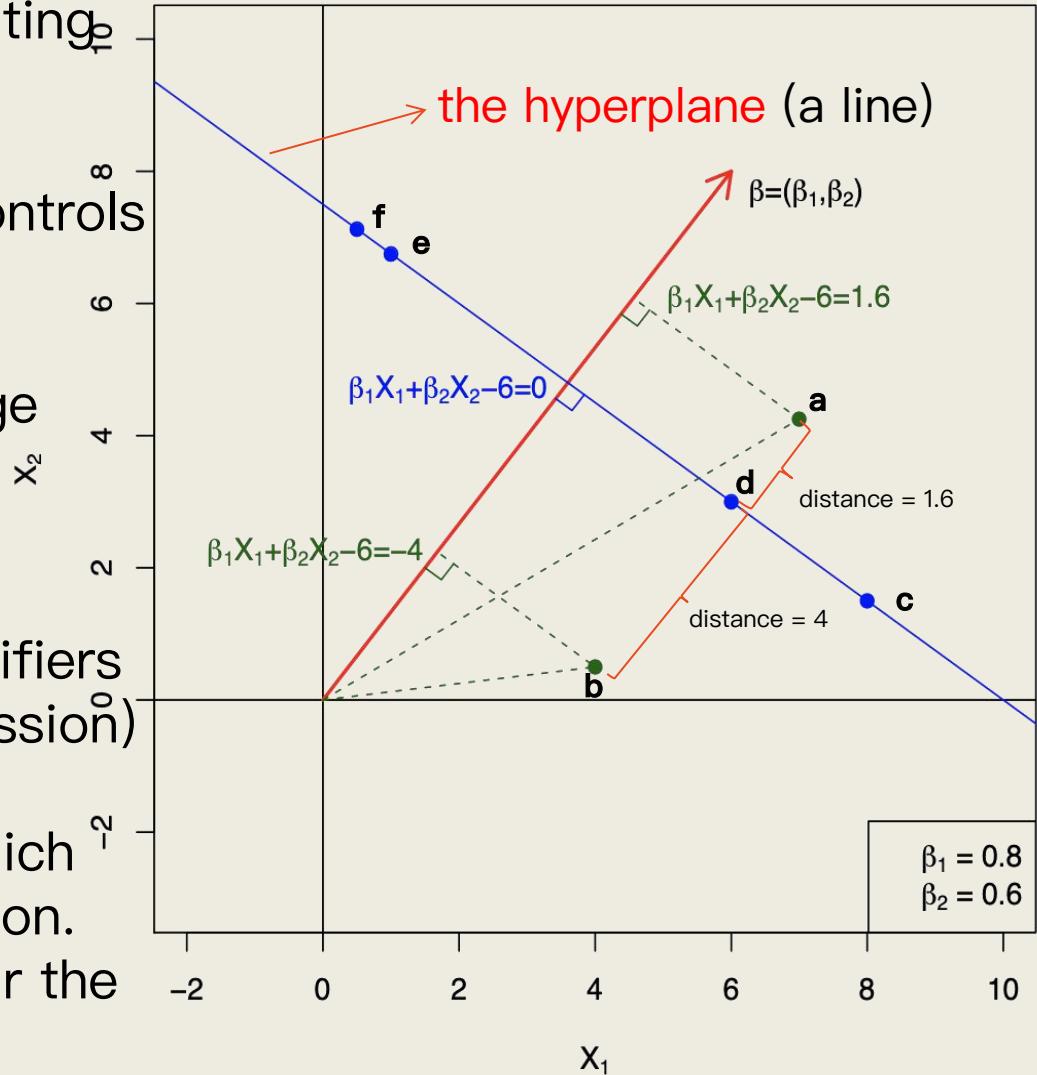
3. **Projections** – green and orange dashed lines.

4. **Connect to classification:**

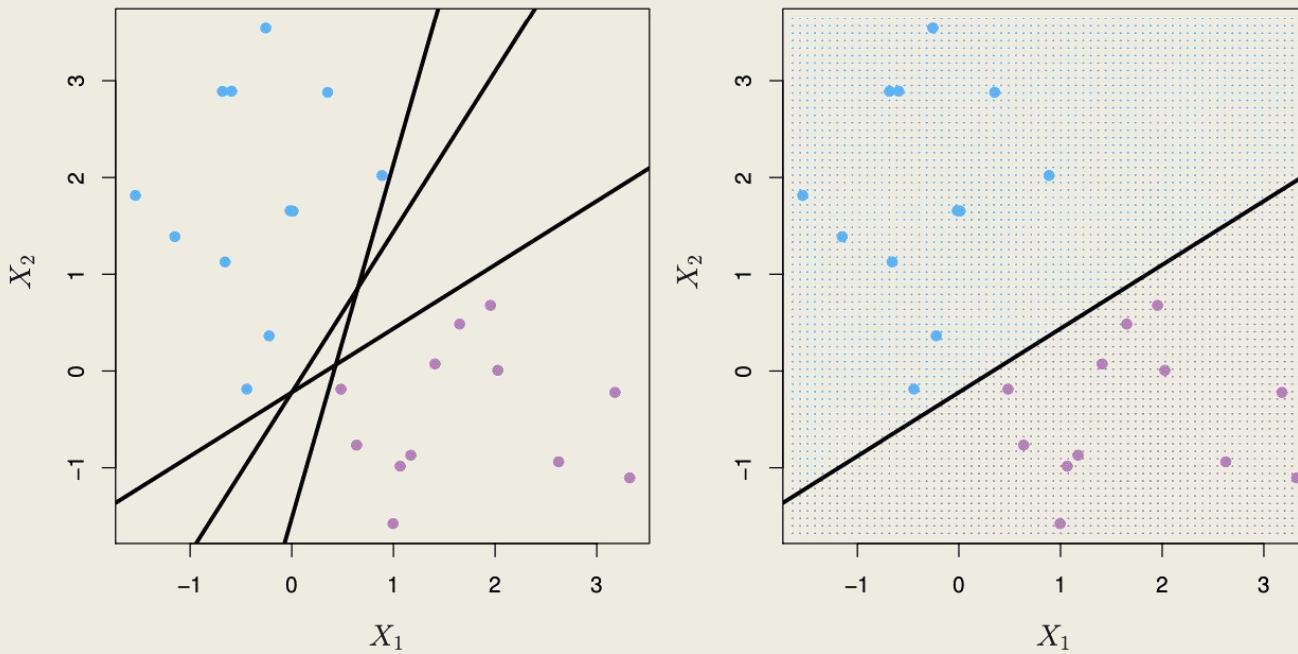
– This is exactly how linear classifiers work (like SVMs or logistic regression)

– **The sign** of $\beta^\top X - 6$ tells you which side of the hyperplane a point is on.

– **The magnitude** tells you how far the point is from the boundary.



Separating Hyperplanes



If $f(X) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$, then $f(X) > 0$ for points on **one side** of the hyperplane, and $f(X) < 0$ for points **on the other**.

If we code the colored points as $Y_i = +1$ for blue, say, and $Y_i = -1$ for mauve, then if $Y_i \cdot f(X_i) > 0$ for all i , $f(X) = 0$ defines a **separating hyperplane**.

Maximal Margin Classifier

- Among all separating hyperplanes, find the one that makes the biggest gap or margin between the two classes.
- Model formulation:

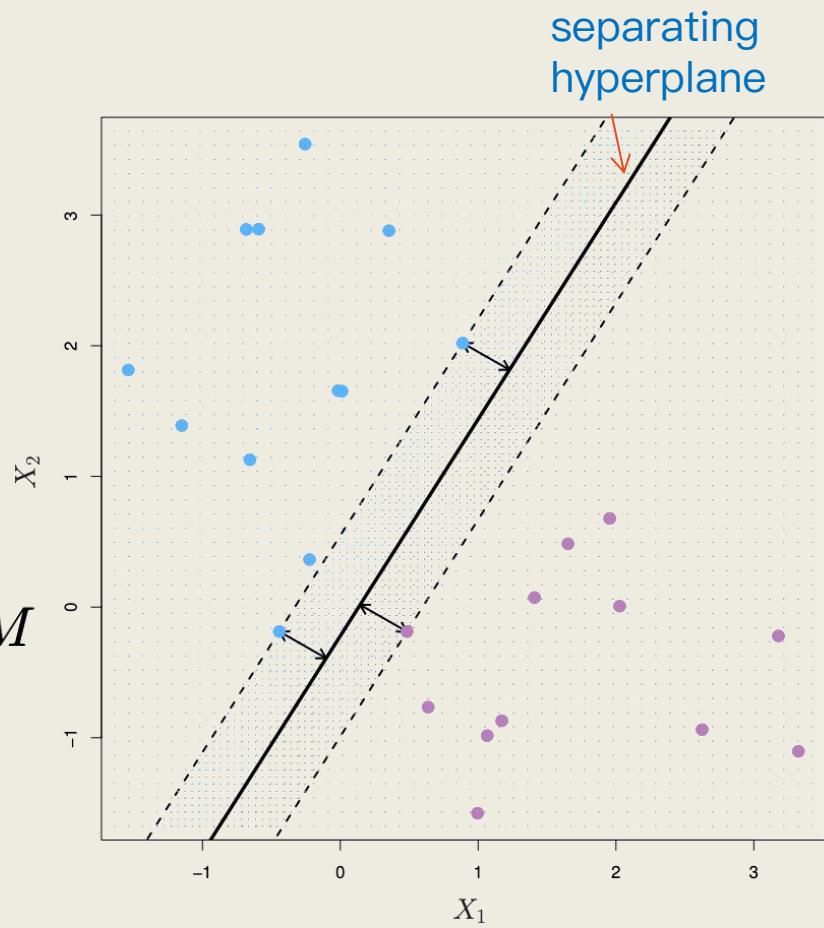
$$\underset{\beta_0, \beta_1, \dots, \beta_p}{\text{maximize}} M \quad \boxed{\text{objective function}}$$

Constraints

$$\begin{cases} \text{subject to } \sum_{j=1}^p \beta_j^2 = 1, \\ y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M \\ \text{for all } i = 1, \dots, N. \end{cases}$$

Inputs $\{\mathbf{x}_i, y_i\}_{i=1}^N$

Outputs β_0, \dots, β_p



Non-separable Data

- The data on the left are not separable by a linear boundary.
- This is often the case,
 - unless the number of data samples N is less than the dimension of inputs p , $N < p$

For example:

- In 3-dimension space, only have two data samples. It's almost always possible to find a plane (hyperplane) that separates the two.
- In high-dimensional spaces with few data points, the points are more "spread out." So, it's easier to find a hyperplane that slices between them without hitting any.



Non-separable Data

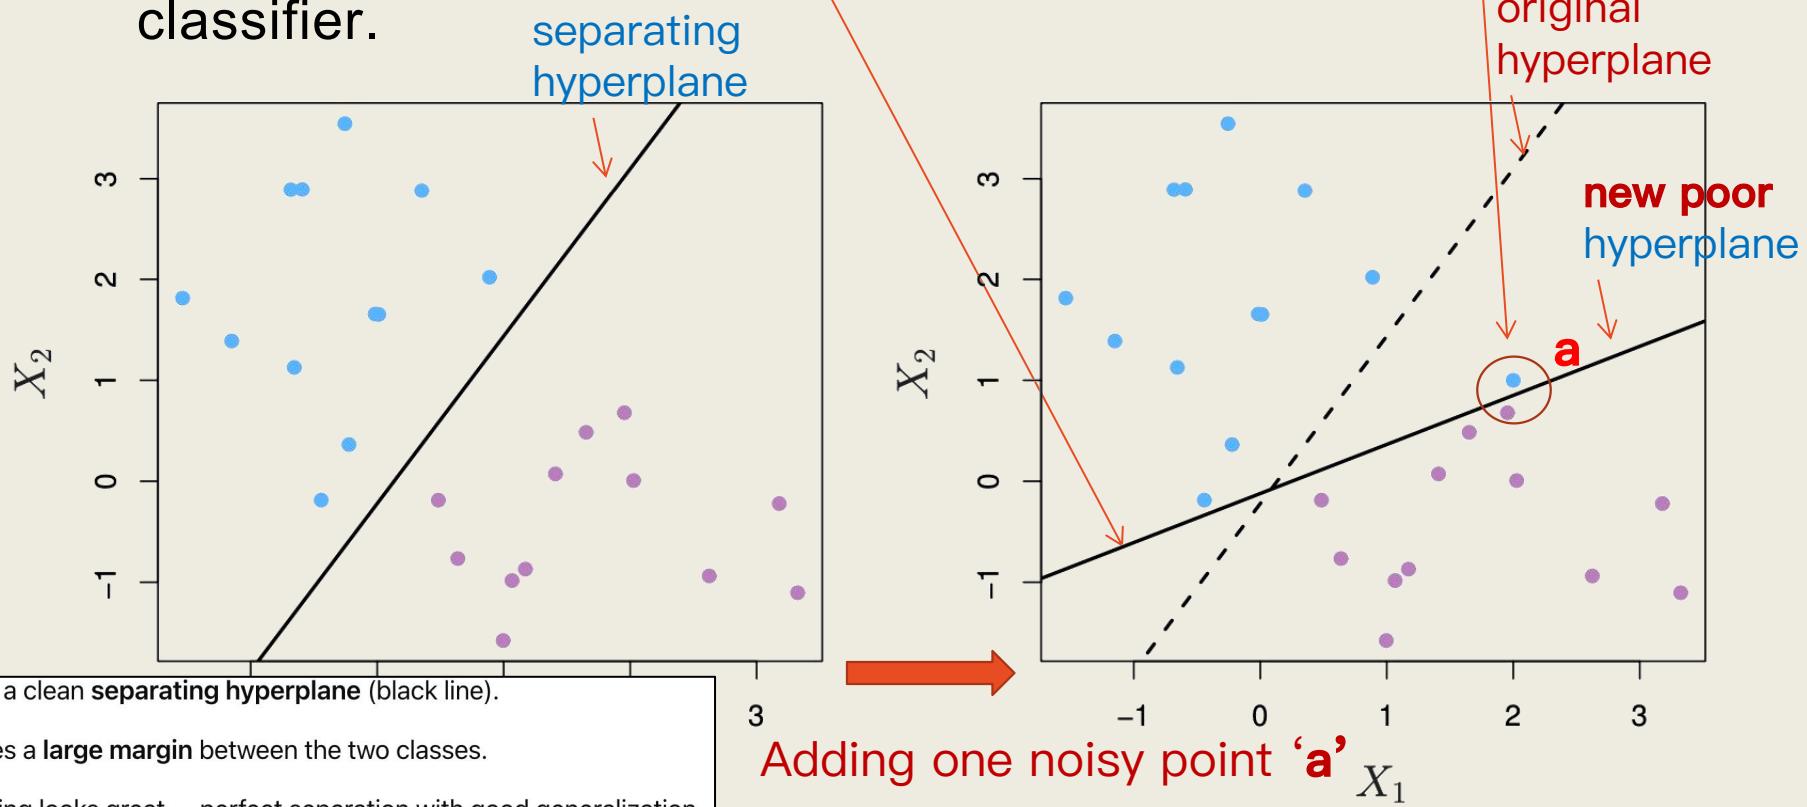
But in practice:

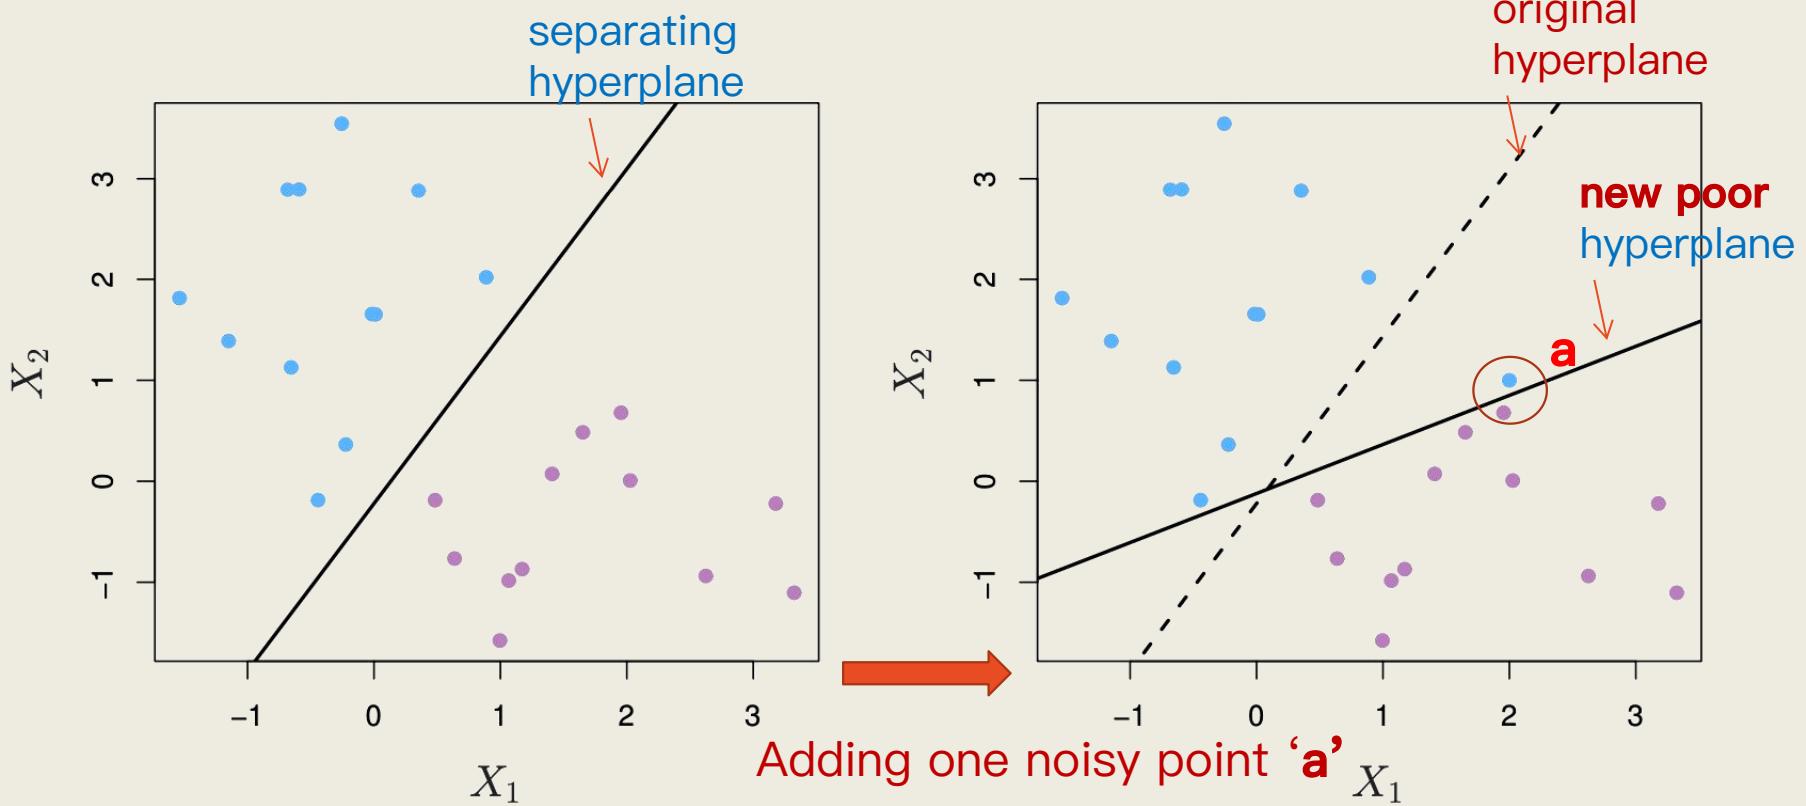
- We usually have $N \gg p$ (lots of data, few features), so:
 - Perfect separation with a linear boundary is rare.
 - That's why we use:
 - Soft margins (allowing some misclassification)
 - Nonlinear transformations (expanding to higher-dimensional spaces)

Noisy Data

- Sometimes the data are separable, but have **noisy samples**.

- This can lead to a **poor solution** for the maximal-margin classifier.





- The new hyperplane (solid line in the right figure) is pulled down and tilted — it now has **a smaller margin**.
- The original hyperplane (dashed line) would've been better if we could just ignore the noisy point.
- The model is now **overfitting to the noise** — giving a poorer solution overall.

Soft Margin

- Just because data are separable, doesn't mean a **hard-margin classifier** will find a good boundary.
- **Noisy samples** can lead to overfitting and poor generalization.
- This motivates the use of **soft margins** — where we allow a few misclassifications in exchange for a better overall boundary.

Support Vector Classifier

- A Support Vector Classifier allows some misclassifications or margin violations (ϵ) by maximizing a **soft margin**, so it can find a better overall boundary in noisy or non-perfectly-separable data

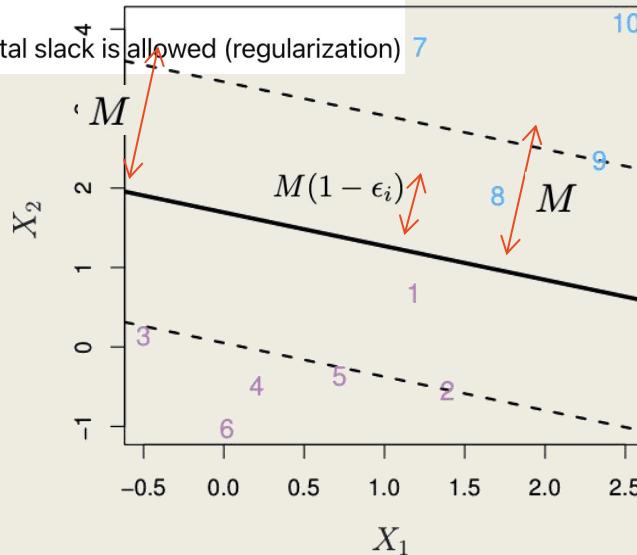
- Model formulation:

M : the margin (distance between support vectors)

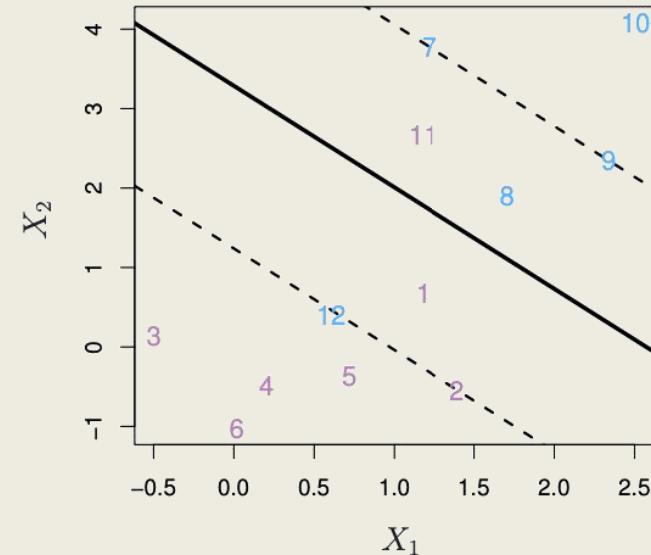
ϵ_i : how much the i -th point violates the margin

- $\epsilon_i = 0$: perfectly classified, outside the margin
- $0 < \epsilon_i < 1$: inside the margin, but on the correct side
- $\epsilon_i > 1$: misclassified

C : controls how much total slack is allowed (regularization)



$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} \quad M \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 = 1, \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i), \\ & \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C, \end{aligned}$$



Hard Margin vs. Soft Margin

- Hard margins require perfect separation — very sensitive to noise.
- Soft margins (Support Vector Classifier) balance:
 - Maximizing margin
 - Minimizing total violations
- Controlled by slack variables ϵ_i and regularization parameter C

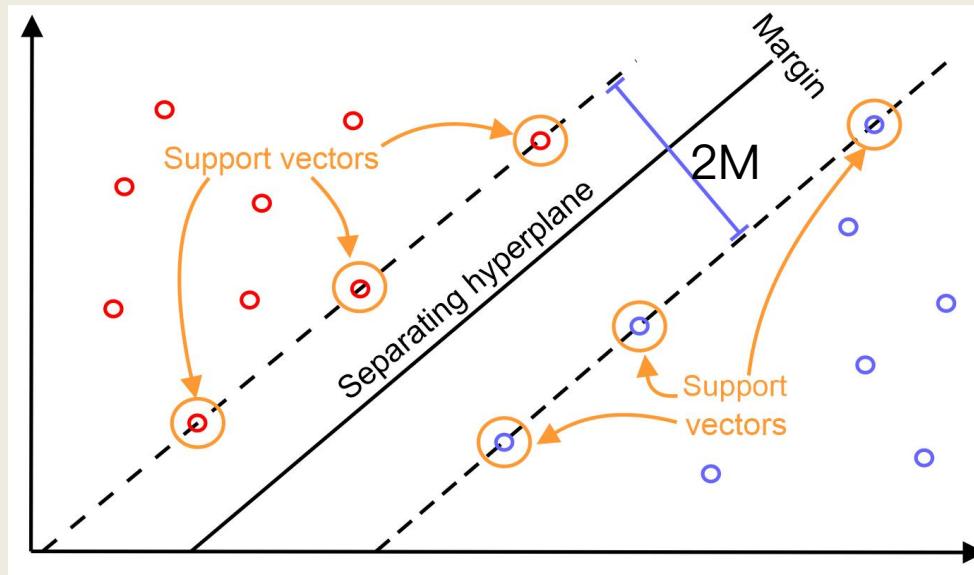
Support Vector Classifier

- **Support vectors:** vectors on two dashed lines

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) = M$$

- **Separating hyperplane (linear boundary):**

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) = 0$$



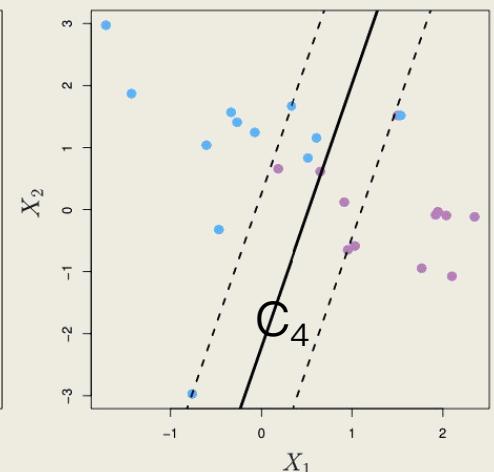
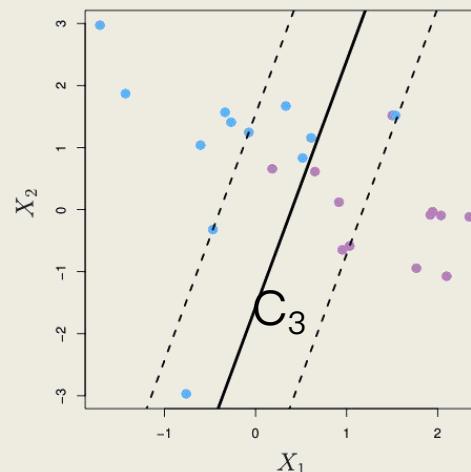
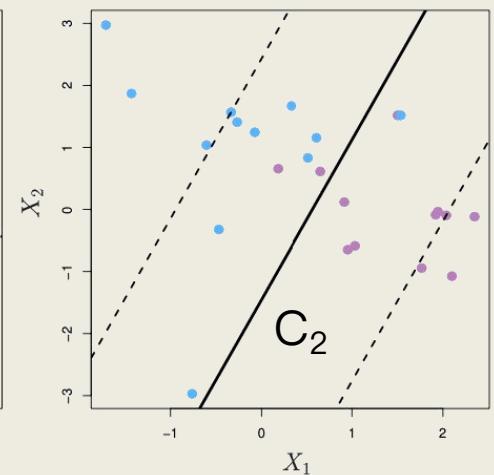
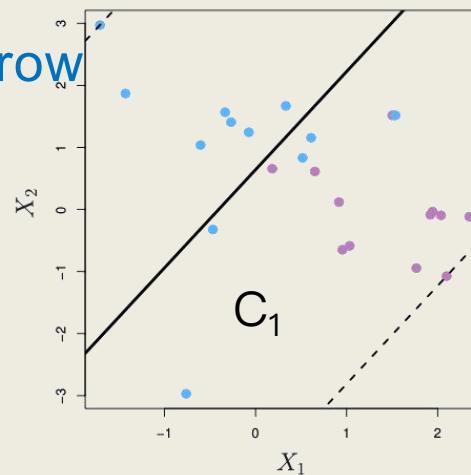
Hyper-parameter C

- where C is a regularization hyper-parameter

- tune(choose) the optimal C by grid search

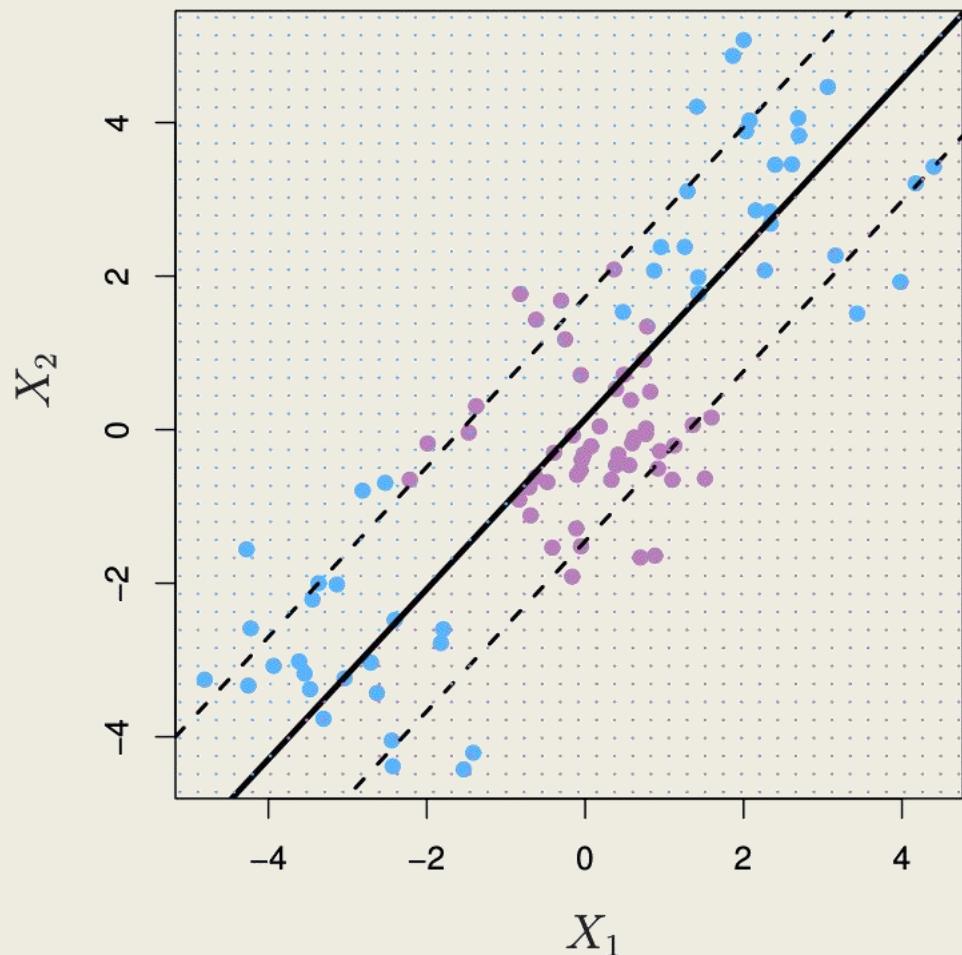
- a **small** C results in **narrow margins**

- as C increases the margins **widen** allowing more violations



Limitation

- Sometime a **linear boundary** simply won't work, no matter what value of C.
 - What shoud we do?
 - Feature Mapping**



Data Transformation – Feature Expansion

- Enlarge the space of features:
 - feature expansion: $X_1^2, X_1^3, X_1X_2, X_1X_2^2, \dots$
 - p -dimensional space is enlarged to a q -dimensional space, where $q > p$.
- Fit a support-vector classifier in the enlarged q -dimensional space.
- This results in **non-linear decision boundaries** in the original space.

Example: Suppose we use $(X_1, X_2, X_1^2, X_2^2, X_1X_2)$ instead of just (X_1, X_2) . Then the decision boundary would be of the form

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 = 0$$

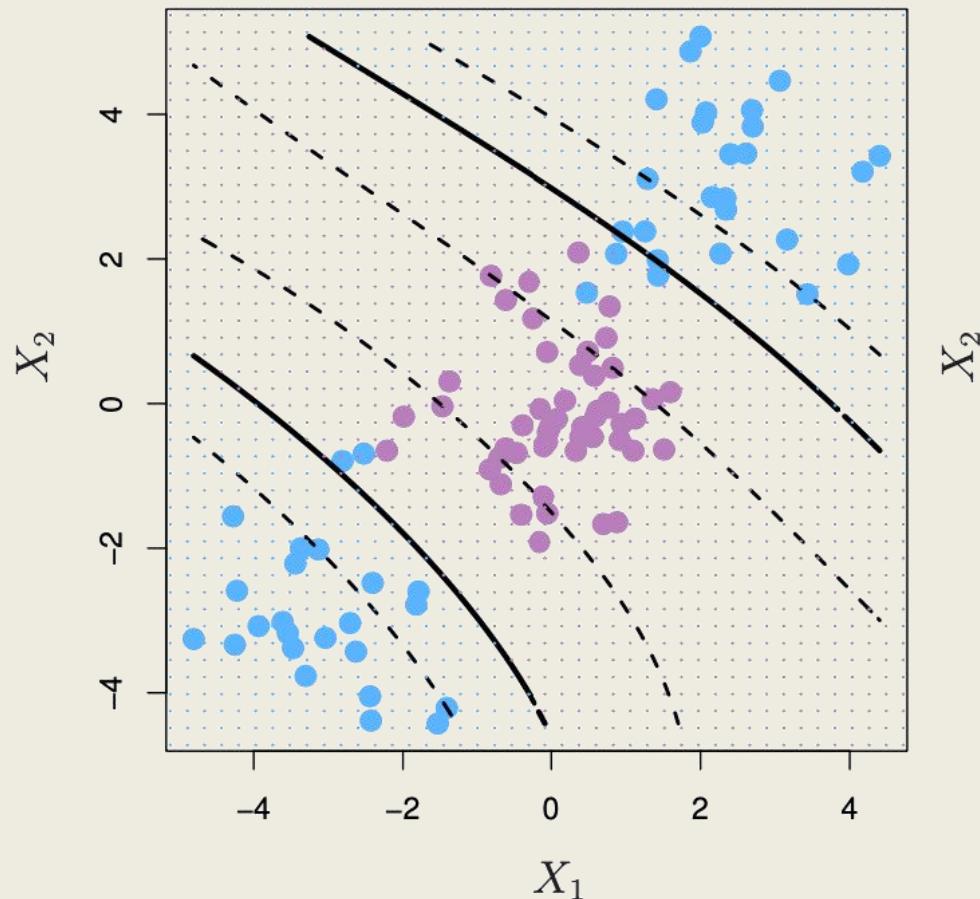
This leads to nonlinear decision boundaries in the original space

Cubic Polynomials

Here we use a basis expansion of cubic polynomials

From 2 variables to 9

The support-vector classifier in the enlarged space solves the problem in the lower-dimensional space



$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 + \beta_6 X_1^3 + \beta_7 X_2^3 + \beta_8 X_1 X_2^2 + \beta_9 X_1^2 X_2 = 0$$

Extend SVMs handling more than 2 classes?

- Standard SVMs are binary classifiers, they only distinguish between two classes. What do we do when we have $K > 2$ classes?

OVA One versus All. Fit K different 2-class SVM classifiers $\hat{f}_k(x)$, $k = 1, \dots, K$; each class versus the rest. Classify x_* to the class for which $\hat{f}_k(x_*)$ is largest.

OVO One versus One. Fit all $\frac{K(K - 1)}{2}$ pairwise classifiers $f_{\ell}(x)$. Classify x to the class that wins the most pairwise competitions.

Which to choose? If K is not too large, use OVO.

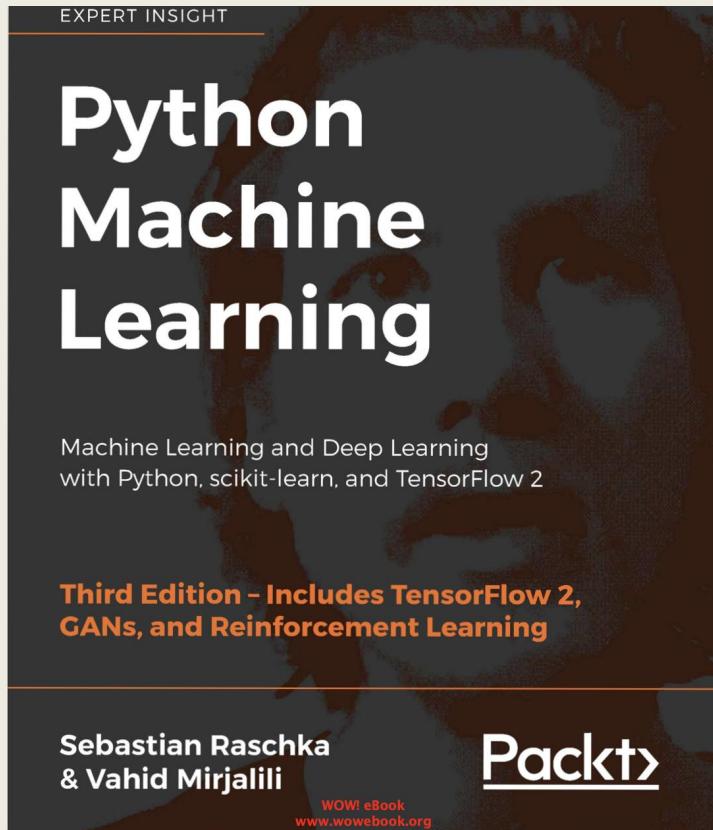
- OVA is simpler and faster, especially when K is large.
- OVO is more precise, but requires more classifiers and computation.

Summary

- **Support Vector Machines**

- Hyperplane
- Maximal Margin Classifier
- Support Vector Classifier
- Feature Expansion

References



- [Understanding Support Vector Machines In Depth Tutorial](#)
- <https://cambiotraining.github.io/intro-machine-learning/svm.html>
- Sebastian Raschka, etl., ‘Python machine learning’, Third Edition